

Scribe Notes for Lecture 14 - 2/15/16

Prepared by Mike Rabbitz, Parth Patel, Alexander Patin, and Shawn Philip

Overview

This lecture is a continuation of lecture 13, where we were talking about the Process Environment and the difference between a program and a process. In lecture 14, we covered topics relating to Process Management. During this lecture, we discussed how processes work and how the OS distinguishes the difference between each of these processes.

The following topics were covered in class:

- Process Identity
- Processes in Unix
- Process Trees in Unix
- Process Lifecycle
- Fork()

I. Process Identity

Every process has its own unique ID, also known as the **pid**. This way the OS can keep track of each process. The number of processes assigned is finite. If we want to find the ID of the current process, we need to use:

```
pid_t getpid(); //Returns the process ID for the current process
```

The `pid_t` data type is a signed integer type which is capable of representing a process ID. The **getpid()** function returns the process ID of the current process.

The functions that are covered in this lecture can be found in these header files' libraries:

```
#include <sys/types.h>
#include <unistd.h>
```

Another way of identifying a process is through the Task Manager. The applications you may see in the Task Manager are actually made up of multiple processes. These processes operate by their communication with each other. On Windows, you can view processes by running “Task Manager”. In Linux/Unix, the **\$top** command displays and updates information about the top cpu processes. It provides a dynamic real-time view of a running system. It can display system summary information, as well as a list of **processes** or threads currently being managed by the **kernel**. The types of system summary information shown, and the types, order, and size of information displayed for tasks, are all user-configurable. Raw cpu percentage is used to rank the processes. If **number** is given, then the top **number** processes will be displayed instead of the default. The first process to run is called “**init**”, which has a pid value of 1.

It is important to note that there are a **finite** number of processes that can be concurring. It is a very large number, but poor handling of processes can cause buildup on the CPU. However, when processes are terminated, their process id can be recycled.

Here’s an example of using the *top* utility:

```
ubuntu: ~
top - 20:35:08 up 2 min,  2 users,  load average: 2.44, 1.31, 0.51
Tasks: 349 total,   3 running, 345 sleeping,   0 stopped,   1 zombie
%Cpu(s): 71.1 us, 28.3 sy,  0.3 ni,  0.0 id,  0.0 wa,  0.3 hi,  0.0 si,  0.0 st
KiB Mem:  2042652 total, 1047436 used,  995216 free,   55088 buffers
KiB Swap: 2094076 total,    0 used, 2094076 free.  380528 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
2398	shawn	20	0	1231196	178096	41048	S	89.1	8.7	0:11.43	compiz
1268	root	20	0	346728	122516	41628	S	3.3	6.0	0:05.27	Xorg
2185	shawn	20	0	376344	5036	3648	S	1.6	0.2	0:00.32	ibus-daemon
2213	shawn	20	0	488016	15708	10744	S	1.3	0.8	0:00.41	ibus-ui-gtk3
2660	shawn	20	0	584220	19908	12896	S	1.0	1.0	0:00.41	gnome-terminal
1694	root	20	0	165504	4684	3776	S	0.7	0.2	0:00.44	vmtoolsd
2441	shawn	20	0	316236	18536	14732	S	0.7	0.9	0:01.00	vmtoolsd
2707	shawn	20	0	29280	1812	1152	R	0.7	0.1	0:00.33	top
2233	shawn	20	0	578488	16960	11580	S	0.3	0.8	0:00.58	unity-panel-ser
2735	shawn	20	0	687420	17844	9652	S	0.3	0.9	0:01.00	unity-scope-hom
2797	shawn	39	19	119568	66188	46764	R	0.3	3.2	0:01.59	apt-check
1	root	20	0	33784	3152	1468	S	0.0	0.2	0:04.61	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.03	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:00.16	ksoftirqd/0
4	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kworker/0:0
5	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/0:0H
6	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kworker/u128:0
7	root	20	0	0	0	0	S	0.0	0.0	0:01.36	rcu_sched
8	root	20	0	0	0	0	R	0.0	0.0	0:00.59	rcuos/0
9	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcuos/1

Figure 1. The top utility helps identify running processes.

II. Processes in Unix

All of the processes in Unix are related with the concept of a parent-child relationship. By associativity, each process is like a grandfather to all of the other ones. The first process started during booting of the computer system is **init** (short for initialization.) It is a special daemon process that is unique and is also known as the ancestor of all processes, whether direct or indirect. Even though it doesn't really do much, it initially gets everything running, and always stays running. It is also good to note down that every process has a parent. In Figure 1, you can see the list of processes and the PIDs associated with each.

III. Process Trees in Unix

Figure 2 is an example of parent-child processes. Not all processes have to be the child of the previous process made. This is because processes are given identity files based on the order in which they are created; hence why in Figure 2, Parent 1 has Children labeled 2, 5, and 6, and not simply 2, 3, and 4. Child 2 went onto be the Parent of Child 3 and Child 4 before the Child 5 of Parent 1 was created. This concludes the fact that other processes can continue to be created from an ancestor process.




Figure 2. Parent-child relationship.

IV. Process Lifecycle

1. First, a process is created.
 - a. The parent process “clones” itself when a new process is created.
2. Next, the process executes a program.
 - a. It loads the program from a file.
 - b. Executes the code afterwards
3. The program exits.
 - a. A parent is occasionally dependant on the child finishing first. In this case, the parent will pause execution until the child has finished.

Example: Suppose a shell is opened, which in itself is a process. In the occurrence of command **top** is the creation of a new **top** process, which is the parent of the shell. The parent then waits for the child to finish, and there is nothing you can do in the shell until the **top** process exits. This is a good example of why a parent would want to wait for the child to finish first.

V. Fork()

Process Creation: Fork

1. The system call for fork is **pid_t fork()**; and it is used to create processes. It takes no arguments and the significance of the return is the ID of the **new process**. The purpose of **fork()** is to create a **new** process, which becomes the **child** process of the caller. After a new child process is created, the result will be **2 processes** running (possibly,) each with its own code. **Both processes** will then execute the next instruction following the **fork()** system call. Therefore, we have to distinguish the parent from the child. This can be done by testing the value returned by **fork()**, which is **different for each process**.

Note: fork() returns twice, once in the parent and once in the child.

There are three cases for the returned value by fork() :

- If **fork()** returns a negative value to the parent, the creation of a child process was unsuccessful (**an error**)
- **fork()** returns a zero to the newly created child process
- **fork()** returns a positive value, the **process ID** of the child process, to the parent. The returned process ID is of type **pid_t** defined in **sys/types.h**. Normally, the process ID is an **integer**. Moreover, a process can use function **getpid()** to retrieve the process ID assigned to this process.

*The child's program counter is initialized to that of the parent. However, because each has its own unique pid, a system call to **fork()** can tell which process is the child. After this, both processes resume from the **same point after fork()**, but with different return values

At the end of the lecture, we were going to learn about **Process Execution**. The next lecture should introduce more of this topic in Process Management.